

Alias	Display name	Hugging Face / served ID	Backend	Role	Base (parent)	NLI	Embedding
foundation-sec-8b	Foundation-Sec-8B	fdtn-ai/Foundation-Sec-8B	huggingface	domain-adapted	llama-3.1-8b	✓ 52.2%	✓ 52.2%
llama-3.1-8b	Llama 3.1 8B (base)	meta-llama/Llama-3.1-8B	huggingface	base	—	✓ 57.9%	✓ 60.9%
lily-cyber-7b	Lily Cyber. 7B v0.2	segolilylabs/Lily-Cybersecurity-7B-v0.2	huggingface	fine-tuned	mistral-7b-instruct-v0.2	✓ 51.5%	✓ 53.7%
mistral-7b-instruct-v0.2	Mistral 7B Instruct v0.2 (base)	mistralai/Mistral-7B-Instruct-v0.2	huggingface	base	—	✓ 44.5%	✓ 46.9%
zysec-7b	ZySec-7B	ZySec-AI/SecurityLLM	huggingface	fine-tuned	zephyr-7b-beta	✓ 27.1%	✓ 36.0%
zephyr-7b-beta	Zephyr 7B Beta (base)	HuggingFaceH4/zephyr-7b-beta	huggingface	base	—	✓ 37.9%	✓ 43.2%
baronllm-v2	BaronLLM v2.0	Trendyol/Qwen3-14B-BaronLLM-v2-Q8	openai (GGUF server)	fine-tuned	qwen3-14b	✓ 26.2%	✓ 30.8%
qwen3-14b	Qwen3 14B (base)	Qwen/Qwen3-14B	huggingface	base	—	✓ 11.9%	✓ 13.9%
trendyol-cyber-70b	Trendyol Cyber. LLM v2 (70B)	Trendyol/Trendyol-Cybersecurity-LLM-v2-70B-Q4_K_M	openai (GGUF server)	fine-tuned	llama-3.3-70b-instruct	✓ 65.1%	✓ 68.5%
llama-3.3-70b-instruct	Llama 3.3 70B Instruct (base)	meta-llama/Llama-3.3-70B-Instruct	huggingface	base	—	—	—
cybertuned	CyBERTuned	s2w-ai/CyBERTuned-SecurityLLM	choice_ranking	encoder (not generative)	roberta-base	✓ 21.6%	✓ 27.0%
roberta-base	RoBERTa Base (base)	roberta-base	choice_ranking	base (encoder)	—	✓ 21.1%	✓ 26.7%

Accuracy = `accuracy` from each `data/eval/{corpus}/{alias}_summary.json`, shown as a percentage and rounded to one decimal.

Summary

- **12** models in `config/eval_models.json`.
- **11** have been run on both corpora (`.jsonl` + `_summary.json` under `data/eval/nli/` and `data/eval/embedding/`).
- **Not evaluated yet:** `llama-3.3-70b-instruct` (parent base for `trendyol-cyber-70b`).

Pairings (base vs cyber-specialized)

Family	Base	Specialized	Best NLI (specialized / base)
Llama 8B	llama-3.1-8b	foundation-sec-8b	52.2% / 57.9%
Mistral 7B	mistral-7b-instruct-v0.2	lily-cyber-7b	51.5% / 44.5%
Zephyr 7B	zephyr-7b-beta	zysec-7b	27.1% / 37.9%
Qwen3 14B	qwen3-14b	baronllm-v2	26.2% / 11.9%
Llama 70B	(no run)	trendyol-cyber-70b	65.1% / —
RoBERTa	roberta-base	cybertuned	21.6% / 21.1%